

# RNNs and LSTM Quiz

## Multiple Choice Questions

- Q1. What is *not* primary benefit of stacking multiple RNN layers (i.e., stacked RNNs)?  
A. Faster training  
B. Lower memory usage  
C. Better learning of hierarchical features  
D. Simpler architecture
- Q2. Which of the following is *not* the main reason RNNs struggle with long-term dependencies?  
A. Overfitting  
B. Vanishing gradients  
C. Lack of non-linearity  
D. Insufficient data
- Q3. What differentiates an LSTM cell from a standard RNN cell?  
A. It uses ReLU instead of tanh  
B. It introduces gates to control the flow of information  
C. It has fewer parameters  
D. It is a convolutional architecture
- Q4. In a standard RNN, which gate is responsible for deciding how much of the past memory to keep?  
A. Output gate  
B. Forget gate  
C. Input gate  
D. Update gate

## Descriptive Questions

- Q5. Why is the forget gate bias in RNNs often initialized to a high value (e.g., 2 or 3)? Explain its effect on long-term dependency learning.
- Q6. Bidirectional CNNs are often used for POS tagging but not machine translation. Explain why, considering alignment and context flow.
- Q7. Designing an RNN model for variable-length legal documents with long dependencies:  
(a) Choose between vanilla RNN or LSTM.  
(b) Stack layers or keep it shallow?  
(c) Make it unidirectional

Justify each choice based on model behavior and task needs.

- Q8. Consider a vanilla CNN with recurrent weight matrix  $W_h$  sequence length  $100$ . Analyze gradient behavior:  
(1) If  $\|W_h\| = 0.9$ : Will gradients vanish or explode? Justify.  
(2) If  $\|W_h\| = 1.2$ : Will gradients vanish or explode? Justify. Suggest an easy fix and explain how it helps.

*Hint: Consider eigenvalue effects on gradient propagation over time.*