RNNs and LSTM Quiz

Multiple Choice Questions

Q1. What is the primary benefit of stacking multiple RNN layers (i.e., stacked RNNs)?

    A. Faster training

    B. Lower memory usage

    C. Better learning of hierarchical features

    D. Simpler architecture

Q2. Which of the following is not main reason RNNs struggle with long-term dependencies?

    A. Overfitting

    B. Vanishing gradients

    C. Lack of non-linearity

    D. Insufficient data

Q3. What differentiates an LSTM cell from a standard RNN cell?

    A. It uses ReLU instead of tanh

    B. It introduces gates to control the flow of information

    C. It has fewer parameters

    D. It is a convolutional architecture

Q4. In a standard LSTM, which gate is responsible for deciding how much of the past memory to keep?

    A. Output gate

    B. Forget gate

    C. Input gate

    D. Update gate

Descriptive Questions

Q5. Why is the forget gate bias in LSTMs often initialized to a high value (e.g., 2 or 3)? Explain its effect or effect term dependency learning.

Q6. Bidirectional RNNs are often used for POS tagging but not machine translation. Explain why, considering input-output alignment and context flow.

Q7. Designing an RNN model for variable-length legal documents with long dependencies:

    (a) Choose between vanilla RNN or LSTM.

    (b) Stack layers or keep it shallow?

    (c) Make it bidirectional?

    Justify each choice based on model behavior and task needs.

Q8. Consider a vanilla RNN with recurrent weight matrix W_h and sequence length 50. Analyze gradient behavior:

    (1) If $\|W_h\| = 0.9$: Will gradients vanish or explode? Justify.

    (2) If $\|W_h\| = 1.2$: Will gradients vanish or explode? Justify. Suggest an easy fix and explain how it helps.

    Hint: Consider eigenvalue effects on gradient propagation over time.